ANKIT BISEN

Data Engineer | Databricks, Azure & AWS Data Platforms

+91-8085717056 | bisenankit1803@gmail.com | LinkedIn | GitHub | Portfolio

SUMMARY

Data Engineer with 3 years of experience in developing and supporting enterprise data integration, ETL, data validation, and reporting workflows. Skilled in Python, PySpark, SQL, Azure Databricks, AWS S3 and Git for sourcing, transforming, validating, and delivering data from multiple enterprise systems. Contributed to building reliable data pipelines and assisting in data warehouse operations while ensuring data quality and performance.

KEY SKILLS

Data Engineering: ETL/ELT, Medallion Architecture, Incremental Ingestion, Data Modeling, Data Warehousing, Scheduled Workflows
Big Data & Processing: Apache Spark, PySpark, Spark SQL, Azure Databricks, Delta Lake, Auto Loader, Structured Streaming
Cloud & Storage: Azure (Databricks, Data Factory, ADLS Gen2, Synapse), AWS (S3, Glue, Lambda, Athena, IAM)
Data Quality & Ops: Data Validation, Data Quality Checks, Data Cleansing, Deduplication, Schema Enforcement, RCA
AI & Automation: LLM-based SQL Generation, Agentic AI (LangGraph, LangChain), Databricks Genie AI, Python Automation, REST APIs
Languages & Tools: Python, SQL, PySpark, Bash, Git, Jira, Databricks SQL Warehouse, Query Performance Tuning

WORK EXPERIENCE

KPIT Technologies Limited *Jul 2026 – Present*
Engineer — Data Engineering & Analytics (Client: Honda Motor Co., Ltd.) *Bengaluru, India*

- Architected a Medallion Lakehouse pattern (Bronze/Silver/Gold) on Databricks, ingesting raw API JSONs into Bronze Delta tables, enforcing schema validation and deduplication in Silver, and modeling aggregated reporting viewports in Gold.
- Designed and automated a data pipeline to extract ticket-level data from Jira via REST APIs, transforming and loading it into Delta tables for tracking cycle time, resolution time, and identifying process bottlenecks across teams.
- Built an LLM-powered SQL AI Agent that converts business prompts into optimized SQL using intelligent model routing and Pydantic-based stateful execution, with SQL safety checks and automated result visualization.
- Configured Databricks Genie AI against the Databricks SQL warehouse, giving leadership a natural-language interface to query business insights without raising individual data-access requests for parties who didn't require standing access.

KPIT Technologies Limited *Aug 2024 – Jun 2026*
Associate Engineer — Data Engineering (Client: Honda Motor Co., Ltd.) *Bengaluru, India*

- Engineered scalable data pipelines on Azure Databricks to process 2–3 million automotive telemetry records/day (CAN/BLF logs), ingesting raw CSV files into ADLS and transforming them into optimized Delta formats for downstream analytics.
- Implemented a Medallion (Bronze–Silver–Gold) architecture and built ELT pipelines using Azure Databricks.
- Built incremental ingestion using Auto Loader (CloudFiles) with Spark Structured Streaming, ensuring scalable and reliable data ingestion.
- Implemented data validation checks to improve data redundancy and data processing by nearly 40%.
- Engineered early-warning telemetry features (thermal anomalies, RPM irregularities, behavioral outliers) enabling proactive identification and prioritization of at-risk vehicles in the fleet.
- Boosted Spark pipeline efficiency by leveraging broadcast joins, caching, and predicate pushdown, cutting processing time by 30% for daily workloads exceeding 2 million records.

KPIT Technologies Limited *Nov 2023 – Jul 2024*
Software Trainee *Bengaluru, India*

- Developed Python automation scripts that replaced manual processes, increasing team productivity by 50%.
- Published a CLI and npm package that strips or comments console statements in JS/TS projects using Babel AST parsing.

INDEPENDENT PROJECTS

ETL Analyst Agent | Python, LangGraph, LangChain, Pandas, REST API, LLM GitHub

- Built an autonomous ETL agent in Python using LangGraph StateGraph with a conditional LLM–tool loop, executing multi-step pipelines from a single natural-language instruction.
- Developed extract-load and transform-load tools to ingest REST API data, normalize nested JSON via Pandas, and write output in CSV, JSON or Parquet.
- Implemented an LLM-driven transformation layer that infers schema from sampled rows and generates and executes Pandas code at runtime, eliminating per-dataset transformation logic.

EDUCATION

Samrat Ashok Technological Institute, Vidisha *Aug 2019 – May 2023*
B.Tech, Electronics & Communication Engineering

CERTIFICATIONS & ACHIEVEMENTS

Certifications: Databricks Fundamentals · Database Management System **Awards:** KPIT High Flyer Award (2x), KPIT High CSAT Award